

StyleCap

Grounded multi-style video captioning for short clips.

Four styles: formal, sarcastic, humorous-tech, humorous-non-tech.

Tags: Gemma, AMD ROCm, AMD Developer Cloud, Fireworks AI, Docker

PROBLEM

Accuracy + style under 10 minutes

LLM judge scores factual accuracy and tone match.

Hallucinated or generic captions score near zero.

~12 clips, 4 styles each, valid JSON, API key baked in image.

Fact sheet first

ffmpeg 6 frames -> Vision JSON fact sheet -> GLM write 2 candidates x 4 styles

Batched judge picks best. Degradation ladder + schema repair.

Graded: Kimi K2.6 (ground) + GLM 5.2 (write/judge) on Fireworks serverless.

End-to-end on the pod

google/gemma-3-12b-it via vLLM + ROCm - ground, write, judge.

8 public clips in 86 seconds. ~40GB VRAM (82%). Evidence: GEMMA_RESULTS.md

Demo video: live rocm-smi + pipeline run.

SUBMIT

Ready for grader pull

ghcr.io/quantumbyte-01/amd-track2-captioner:v34

Public GHCR | linux/amd64 | parallel workers | CANDIDATES=2

StyleCap | github.com/QuantumByte-01/amd-track2-captioner